


Correct current facts can still leave downstream answers dependent on source history.

A Matched-history design

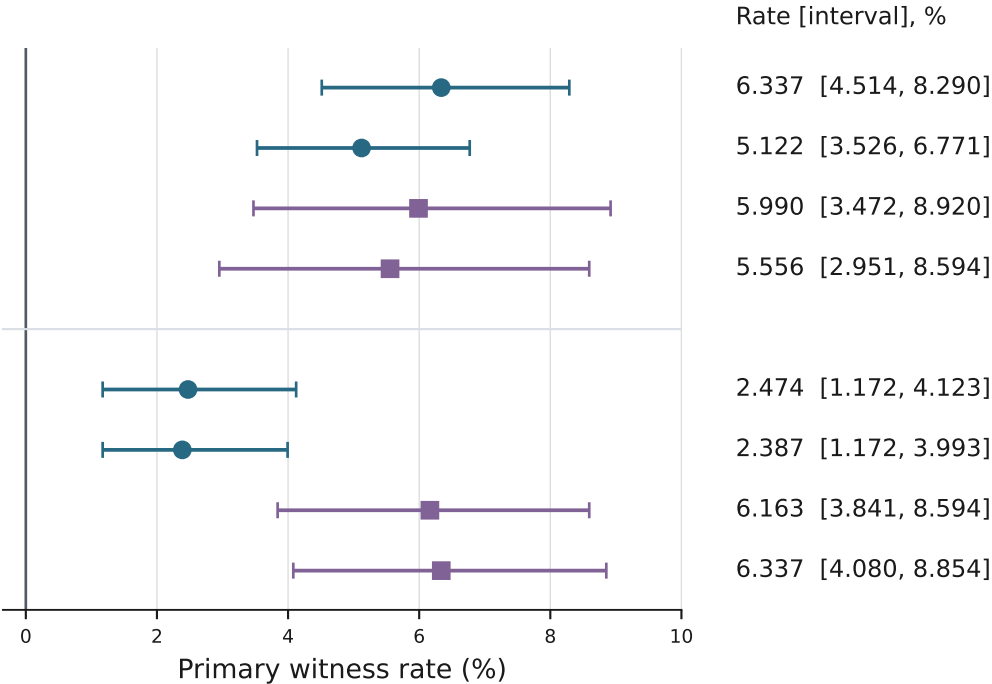


Behavioral test: identical intended facts do not imply identical hidden states or erased history.

B V10 prospective confirmation

Saved primary witness rates and multiplicity-adjusted 99.375% root-bootstrap intervals

Gemma	Constrained learned editor
	Unrestricted consistency-trained editor
	Existing textual correction
	Latest-value wording
Qwen	Constrained learned editor
	Unrestricted consistency-trained editor
	Existing textual correction
	Latest-value wording



All eight adjusted intervals are above zero.

Each cell: 64 roots × 18 fixed joint-question opportunities = 1,152 per seed or textual condition.
Ineligible opportunities remain in the denominator; two learned seeds are averaged within each root.

Witness: correct direct operands in every history + correct native/no-op references + differing joint answers.
No pooling across models or update conditions. Not root prevalence, an answer-error rate, or a method-ranking test.