



A Change vs. factual-update baseline



B Absolute reliability



● Seed mean, 98.75% interval ▼ Seed 0 ▲ Seed 1

64 benchmark cases per model; success requires all 34 questions correct from all four histories. Models are not pooled.